

Tutorial using BEAST v2.x

Structured birth death model

Denise Kühnert

Population structure using the multi-type birth-death model

1 Introduction

In this tutorial (adapted from Tim Vaughan’s [Structured Coalescent tutorial](#)), we will use the [BEAST 2](#) package [bdmm](#) to perform a Bayesian phylogenetic analysis of an influenza data set using the multi-type birth-death model. (Note that both, the structured coalescent and the multi-type birth-death model, are tree priors implemented in BEAST2. Both of them utilize the multi-type tree structure of the MultiTypeTree package. While the structured coalescent is part of the MultiTypeTree package, the multi-type birth-death model has its own package bdmm (aka birth-death migration model).)

The data set used in this tutorial is a thinned (60 sequence) subset of the (980 sequence) data set used in the publication ([Vaughan et al. 2014](#)), which in turn was assembled from publicly-available data sets provided by various authors on [GenBank](#).

1.1 Software Requirements

In order to proceed, ensure you have the latest version (currently 2.4) of [BEAST 2](#) installed. To analyze the inference results you’ll also need a recent version of [Tracer](#) and an up-to-date version of [Google Chrome](#) or [Mozilla Firefox](#).

1.2 Installing the bdmm package

You can easily install this package via BEAUti’s package manager. To do this, follow these steps:

1. Start BEAUti
2. From the *File* menu, select *Manage packages*.
3. Find “bdmm” in the list of packages shown, select it and then click “Install/Upgrade”:

(Note the actual version of bdmm may differ from the version shown in the figure. This is normal.)

Finally, **restart BEAUti**. This is very important. Strange behaviour may result if you do not restart the program.

2 Setting up the analysis using BEAUti

2.1 Loading the Template

A BEAUTI template defines the basic structure and contents of your XML file. Because by default BEAUTI will construct an XML file with standard BEAST trees, rather than MultiTypeTrees, we cannot use the default template (Standard.xml). Hence, from the *File* menu, select *Template* and then choose *MultiTypeBirthDeath*.

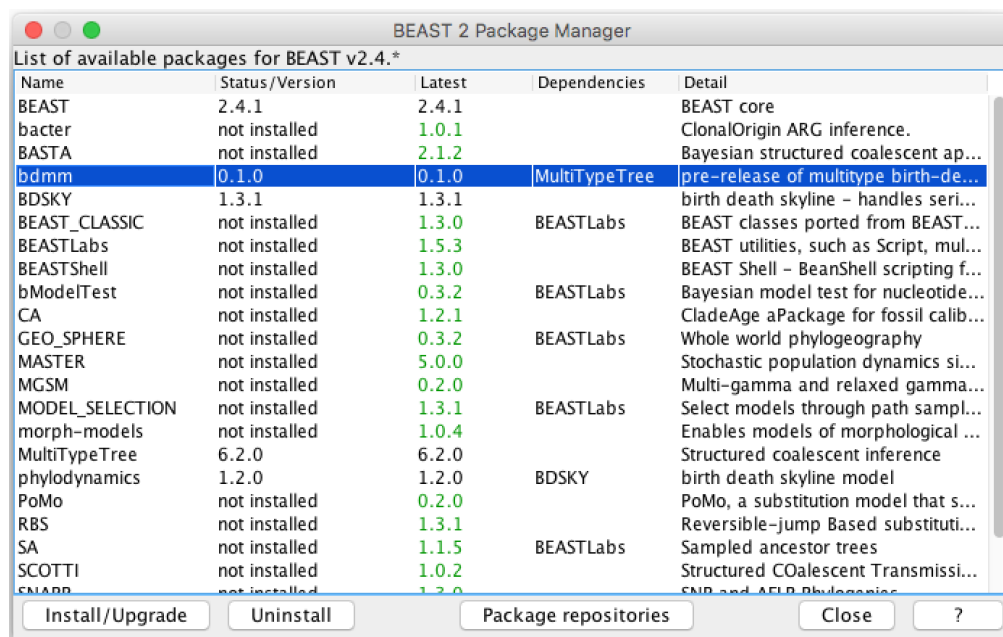


Figure 1: Install bdmm.

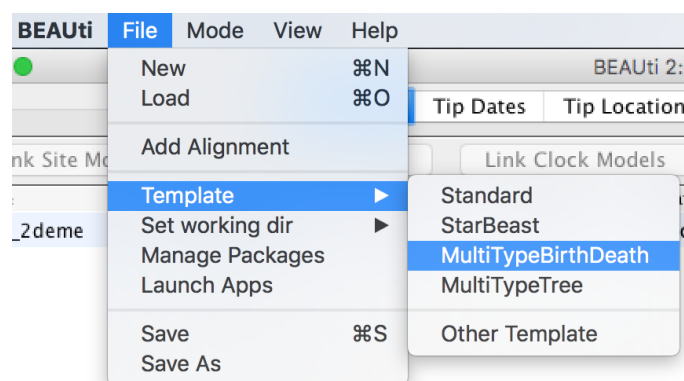


Figure 2: Load the MultiTypeBirthDeath template.

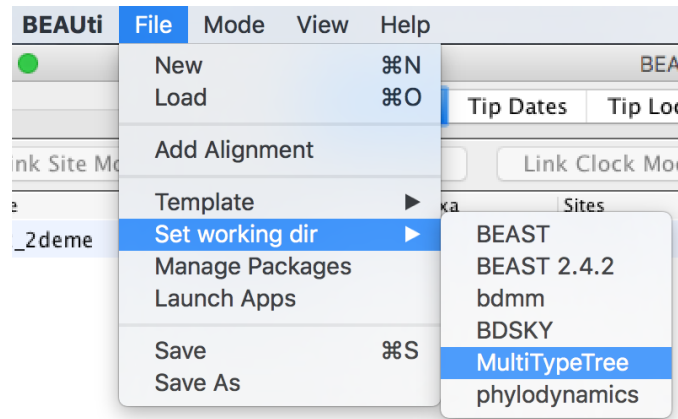


Figure 3: Set the working directory to MultiTypeTree.

2.2 Loading the data

Once the template is loaded, we can load in our example sequence data. In our case, this data is stored in a FASTA file, the first few lines of which look like this:

```
> EU856841_HongKong_2005.34246575
-----GGGATAATTCTATTAACCATGAAGACTATCATTGCTTTGAGCTACATTT
> EU856989_HongKong_2002.58356164
--CAAAAGCAGGGGATAATTCTATTAACCATGAAGACTATCATTGCTTTGAGCTACATTT ...
> CY039495_HongKong_2004.5890411
-----TTCTATTAACCATGAAGACTATCATTGCTTTGAGCTACATTC ...
> EU856853_HongKong_2001.17808219
-----TATTAACCATGAAGACTATCATTGCTTTGAGCTACATTC ...
> CY010084_NewZealand_2005.62739726
-----TATTAACCATGAAGACTATCATTGCTTTGAGCTACATTC ...
> CY007387_NewZealand_2004.63287671
-----TATTAACCATGAAGACTATCATTGCTTTGAGCTACATTC ...
> CY012432_NewZealand_2000.81643836
-----CCATGAAGACTATCATTGCTTTGAGCTACATTT ...
```

The lines beginning with “>” are labels for the sequences immediately following. In general, these labels have no special format, but in this file each label is an underscore-delimited triple. The first element of each triple is the GenBank accession number of the sequence, the second is the geographical region from which it was sampled, and the third is the time at which it was sampled measured in calendar years or fractions thereof. (The ellipses are not in the file, but are used here to indicate that sequence has been truncated.)

In this tutorial we will be using the influenza sequence data which is distributed with MultiTypeTree. To make this easy to find, open the *File* menu and select *Set working dir*. Then, from the submenu that appears, select *MultiTypeTree*.

(This step is not required when loading your own data.)

To load the file, open the *File* menu and select *Add alignment*. This will open a file selection dialog box. The example influenza sequence data file is named `h3n2_2deme.fna`. Assuming you have followed the previous step to set the working directory, this can be found in the `examples/` directory shown when the file selection dialog box loads.




Figure 4: The alignment loaded into BEAUti.

Once this file is loaded, your BEAUti screen should look something like the following:

2.3 Setting up dates

Once the data is loaded, the next step is to specify the times at which the sequences were sampled:

1. Select the “Tip Dates” panel.
2. Check the “Use tip dates” checkbox.
3. Click the “Guess” button at the top-right of the panel. This opens a dialog that allows sample times to be loaded from a file or inferred (guessed) from the sequence labels.
4. Because the times are included as the last element of the underscore-delimited sequence names, choose the “use everything” radio button and select “after last” from the drop-down menu. (Note that the underscore character is already chosen as the delimiter.)

After clicking “OK” you should find that the tip date table is populated with times that match those in the sequence headers, and that the last column of the table contains “heights” (times before most recent sample) calculated from the times:

2.4 Setting up locations

Now that we’ve specified the sampling times, we move on to specifying the sampling locations. To do this, we follow a very similar set of steps to those we used to set the sample times:

1. Select the “Tip Locations” panel. You’ll find that the locations are already populated with default values.
2. Click the “Guess” button at the top-right of the panel. This opens the same dialog that we saw in the previous section.
3. Because the locations are included as the second element of the underscore-delimited sequence names, choose the “split on character” radio button and select group 2 from the drop-down menu. (Note again that the underscore character is already chosen as the delimiter.)

After clicking “OK” you should find that the tip location table is populated with locations that match those in the sequence headers, as follows:

2.5 Substitution model

For this analysis, we will use the HKY substitution model with 4 gamma categories and estimated base frequencies. To configure this in BEAUti, switch to the “Site Model” panel, change the number of gamma categories and select HKY from the drop-down menu (the default option is JC69). We also want the shape and proportionInvariant parameters to be nonzero and estimated to account for heterogeneity between

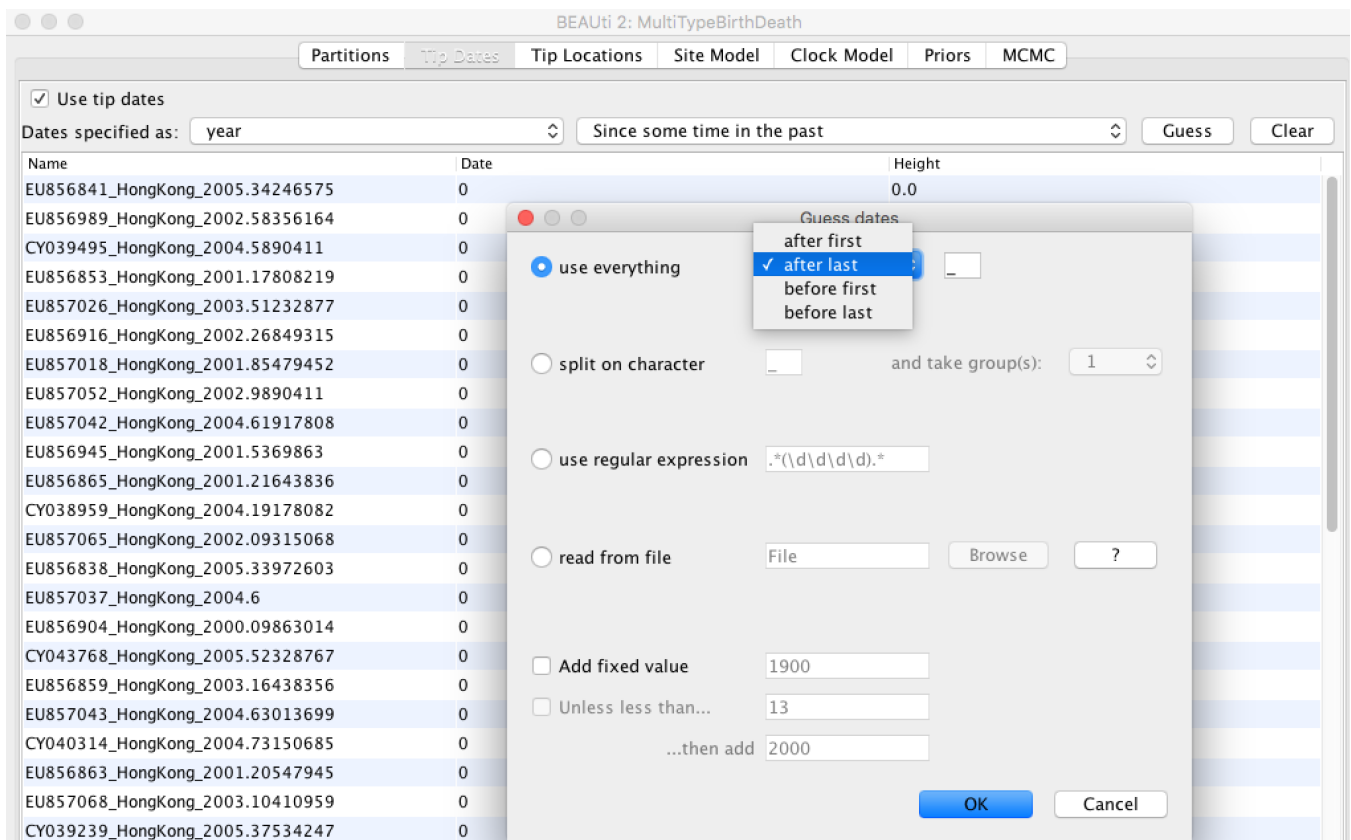


Figure 5: Guessing the tip-dates.

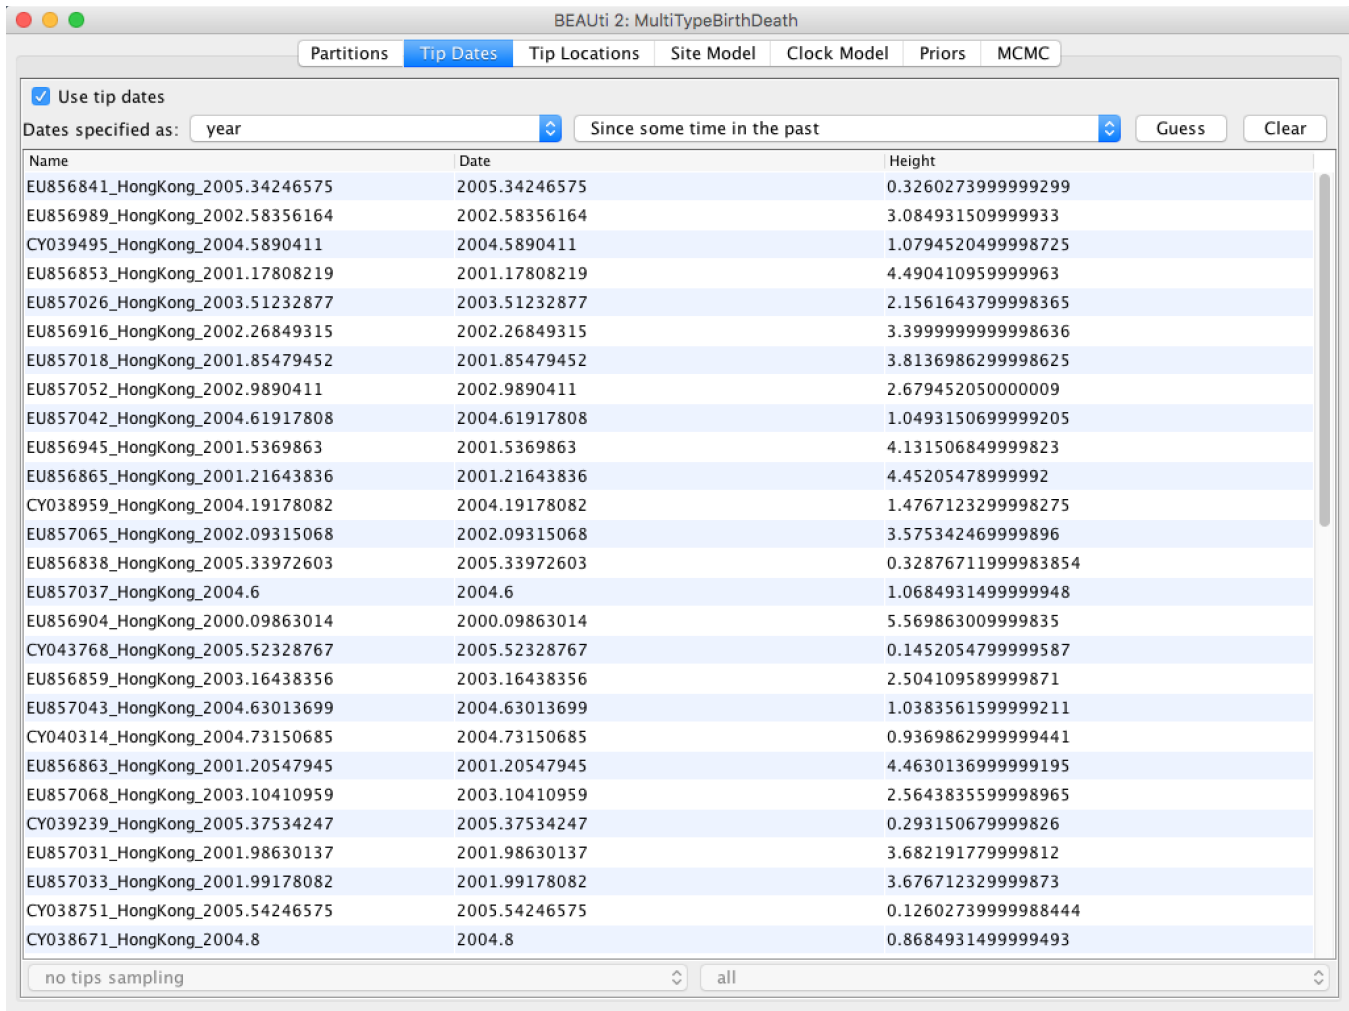


Figure 6: Sampling dates in BEAUti.

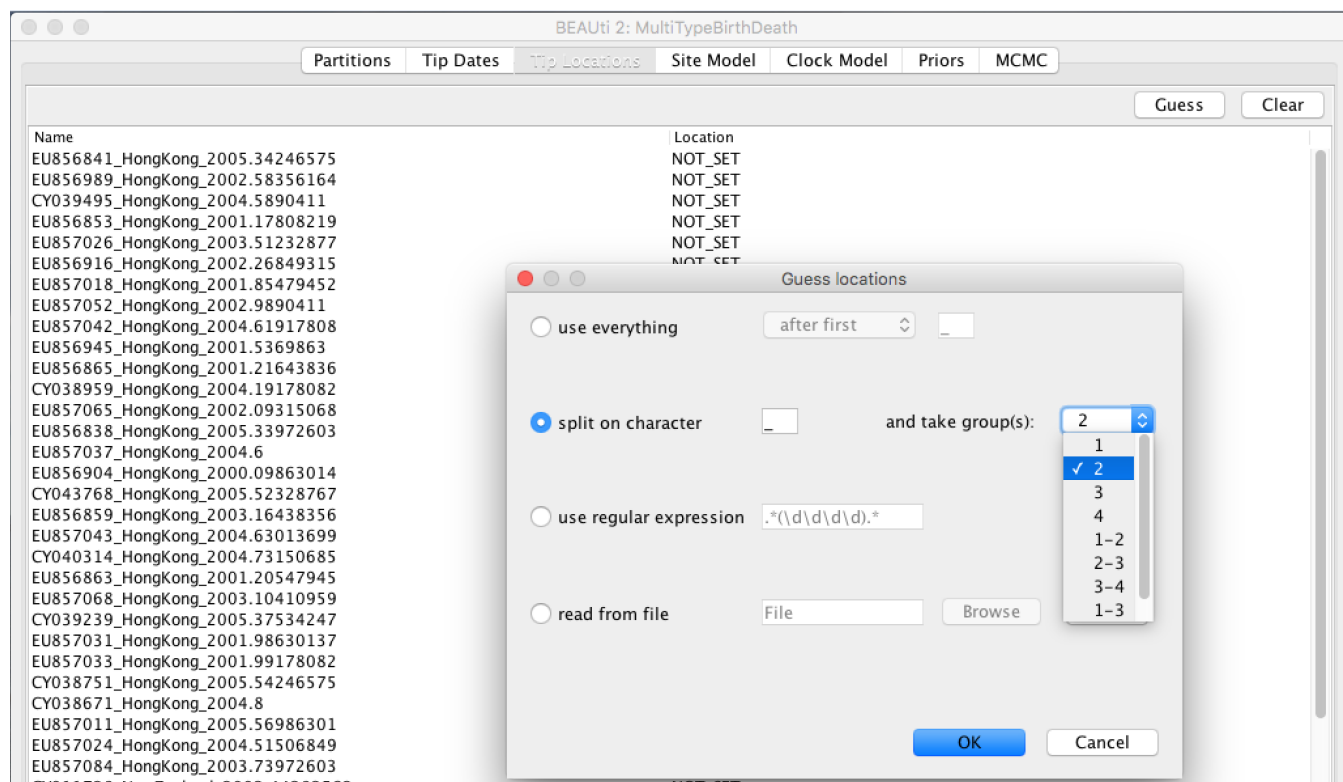


Figure 7: Guessing the locations.

sites in our alignment.

The BEAUi panel should now look like the following:

Note that the “Substitution rate” defined on this panel should be left non-estimated - we use the “Clock rate” defined in the “Clock Model” panel to determine the average per unit time rate of sequence evolution. Used this way, the “substitution rate” is therefore not actually a rate (it’s actually dimensionless) but is instead a rate multiplier that in our case we fix at 1.

2.6 Defining Clock model

For this analysis we assume a strict clock. Since our alignment contains sequences sampled at different times and those times are measured in years, we must use a real clock rate expressed in units of expected substitutions per site per year. Usually the precise value is unknown and so the default behaviour of BEAUi is to assume this rate is to be estimated. We set the clock rate to 0.005, which we know is much closer to the truth than the default 1, to speed up mixing. The Clock Model panel should now look like this:

2.7 Adjusting Priors

Because bdmm is a model-based prior on the (multi-type) tree distribution, setting up the “Priors” panel is a particularly important part of setting up this analysis.

It is important to change the time at which sampling started, counting from the time of the last sample. In




Figure 8: The locations in BEAUi.




Figure 9: Setup of the site model.




Figure 10: Fix the clock rate to speed up mixing.

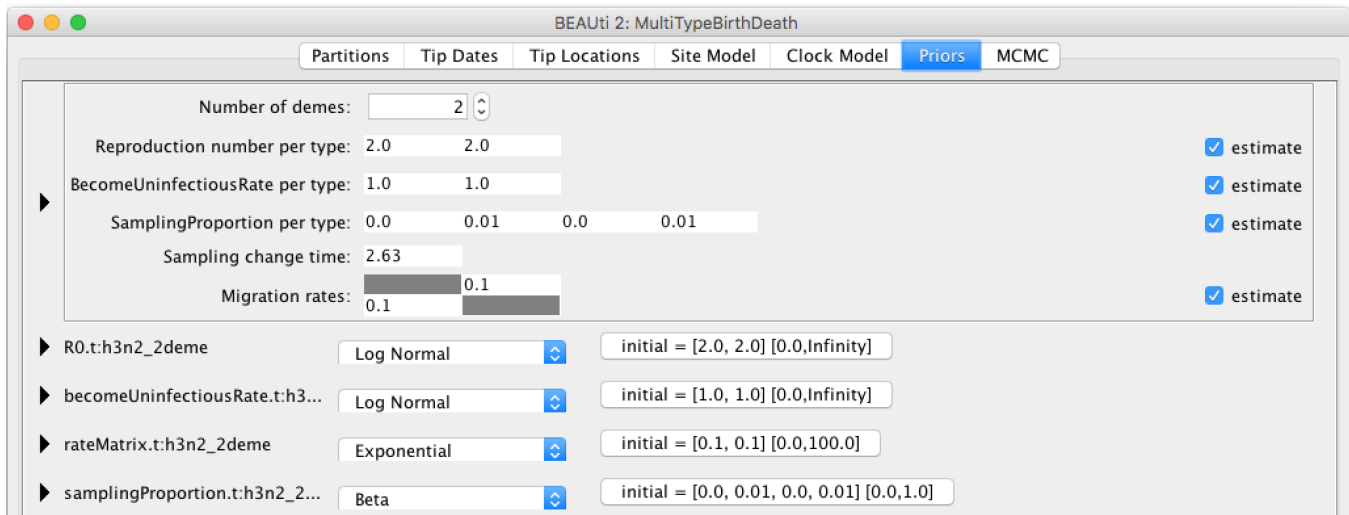


Figure 11: Set the change time for the sampling proportion so it is zero before the time of the first sample.

our case, the time between the first and last sample is 5.57 (rounded UP). Note, that the sampling proportion has 4 entries, let's call them $[s_{11}, s_{12}, s_{21}, s_{22}]$. The values s_{11} and s_{12} belong to type 1 (HongKong) and s_{21} , s_{22} belong to type 2 (New Zealand). The values s_{11} and s_{21} belong to the first interval, which is the time before the first sample, which is why $s_{11}=s_{21}=0$.

When you expand the tree prior element, you can change the condition on survival setting. We'll leave the box checked.

Ensure the value shown in the "State Number" spinner is equal to the number of types present in your model. In this case, the default value of 2 is correct, but in general you should check your data and change this value accordingly.

Click the arrows on the left-hand side of each parameter to alter the details of these priors. For example, we will set the rateMatrix prior distribution to $\text{Exp}(1)$:

We will use the default set-up for the MCMC and save our file as usual.

3 Running the analysis using BEAST

To run the analysis, simply start BEAST 2 in the manner appropriate for your platform, then select the file you generated in the last section as the input file. (Refer to the documentation provided at www.beast2.org for detailed instructions.)

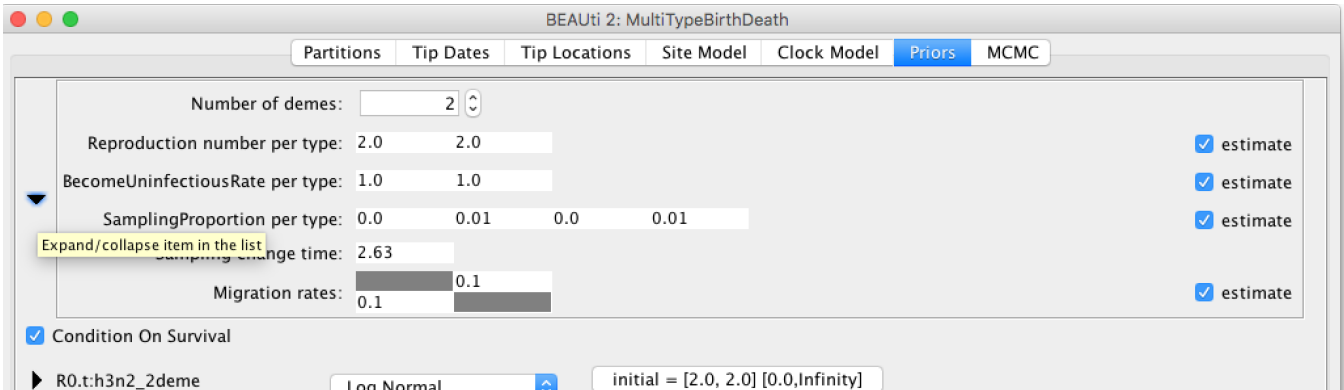


Figure 12: Condition on survival.

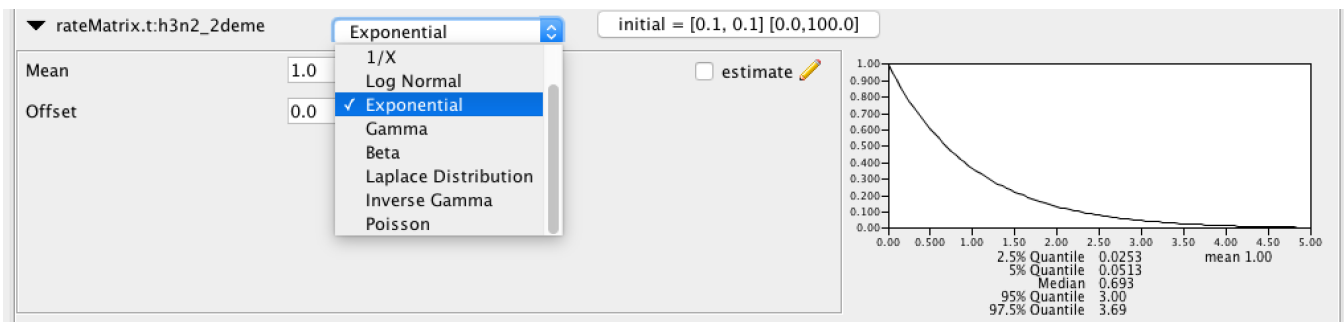


Figure 13: Set the prior for the rate matrix.

4 Analyzing the results

The results of the analysis primarily consist of two parts:

1. The parameter log, which is written to the file `h3n2-bdmm-v2-samplingPrior.log`.
2. The tree log, which is written to `h3n2-bdmm-v2-samplingPrior.h3n2_2deme.trees`.

In addition, the file `h3n2-bdmm-v2-samplingPrior.h3n2_2deme.map.trees` contains the running estimate of the MAP tree as a function of MCMC step number, while the file `h3n2-bdmm-v2-samplingPrior.h3n2_2deme.typedNode.trees` is the TreeAnnotator-compatible file we'll use to assemble a summary tree.

4.1 Parameter log file analysis

We can use the program [Tracer](#) to view the parameter log file. To do this, start Tracer and then press the “+” button in the top-left hand corner of the window (under “Trace files”). Select the log file for this analysis (`h3n2_2deme.log`) from the file selection dialog box. The “Traces” table will then be populated with parameters and summary statistics corresponding to our multitype birth-death analysis.

Important traces are: * `R0.t:h3n2_2deme1` and `R0.t:h3n2_2deme2`: These give the effective reproduction numbers for deme 1 (Hongkong) and 2 (New Zealand), respectively.

- `rateMatrix.t:h3n2_2deme1` `rateMatrix.t:h3n2_2deme2`: These give the (per lineage per year) migration rates from deme 1 to 2 and vice versa.

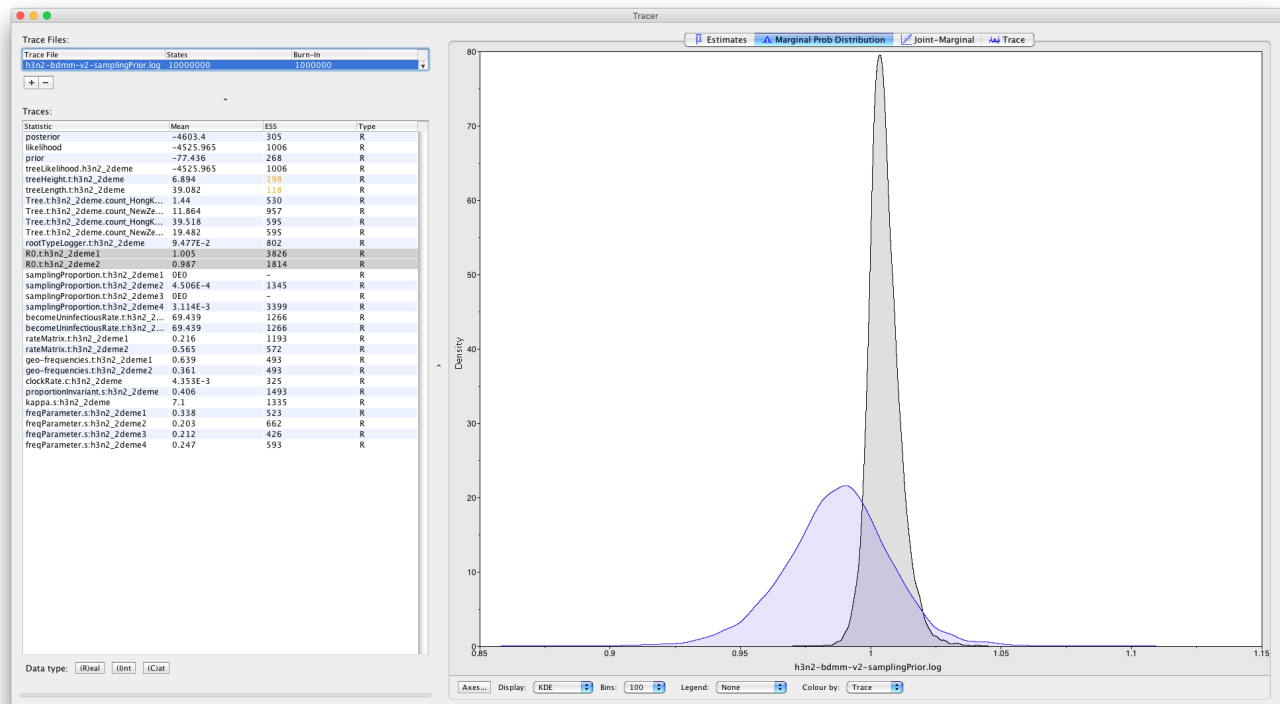


Figure 14: Estimated ‘ R_0 ’ marginal posteriors.

- **Tree.t:h3n2_2deme.count_HongKong_to_NewZealand:** these give the actual number of ancestral migrations from HongKong to New Zealand.

The panels tabs at the top-right of the window can be used to display one or more selected traces in various ways. For example, selecting the two R0 traces and choosing the “Marginal prob distribution” panel results in the following useful comparison between the sampled population size marginal posterior distributions:

Note that some of the ESS values are still less than 200 - the arbitrary threshold for acceptability. If this analysis were part of a serious study you would want to run the chain for another few million iterations to improve these. (In BEAST 2, analyses can be resumed - the samples you’ve already acquired will not be wasted.) For the purposes of this tutorial, however, these values are acceptable.

4.2 Tree log visualization

The popular phylogenetic tree visualizer [FigTree](#) can be used to visualize the sampled trees. Be warned, however, that FigTree currently takes an extremely long time to load even relatively small (a few megabyte) MultiTypeTree logs.

For this reasons we suggest using [IcyTree](#) to view tree log files and maybe switching to FigTree to visualize summary trees as discussed in the next section. (Also, IcyTree can be used to export individual trees from a large log file for subsequent viewing using FigTree.) IcyTree is a tree viewer that runs in a web browser. It runs best under recent versions of [Google Chrome](#) and [Mozilla Firefox](#) (in that order).

To view MultiTypeTree log files using IcyTree, simply navigate to the IcyTree web page, select “Load from

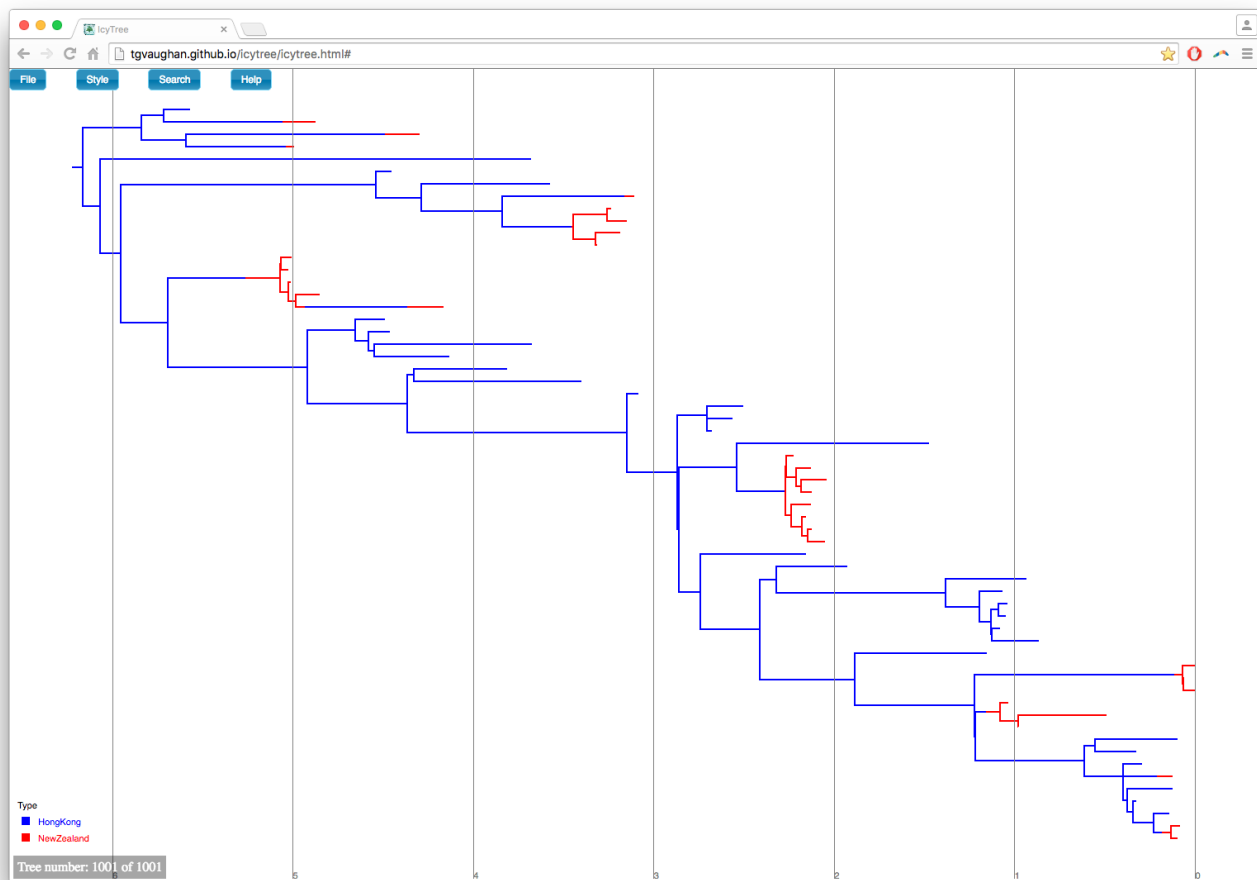


Figure 15: The MAP multi-type tree in IcyTree.

file” from the “File” menu, then select one of the tree log files using the file selection dialog. Once the file is loaded you will see the first tree it contains. In order to select a different tree, move the mouse pointer over the box in the lower-left corner of the window. This box will expand to a small dialog containing buttons allowing you to navigate between trees. The ‘<’ and ‘>’ buttons move in steps of 1 tree, while ‘<<’ and ‘>>’ move 10% of the tree file per click. You can also directly enter the index of a tree. (Note that there are keyboard shortcuts for almost all commands in IcyTree and that these can be found by selecting “Keyboard shortcuts” from the “Help” menu.)

Initially the trees edges will be uncoloured. To colour the edges according to the edge type, open the “Style” menu, navigate to the “Colour edges by” submenu and select “type”. A legend and axis can be added by choosing “Display legend” and “Axis > Age” from the same menu.

The following shows the final tree of `h3n2-bdmm-v2-samplingPrior.h3n2_2deme.map.trees` in IcyTree, which represents our sampled estimate of the MAP multi-type tree:

While IcyTree is useful for rapidly visualizing the results of an analysis, it is not nearly as feature-rich as FigTree and not as capable for producing publication-quality graphics. Happily, however, IcyTree can extract single trees from larger log files. Simply navigate to the desired tree, open the “File” menu, choose

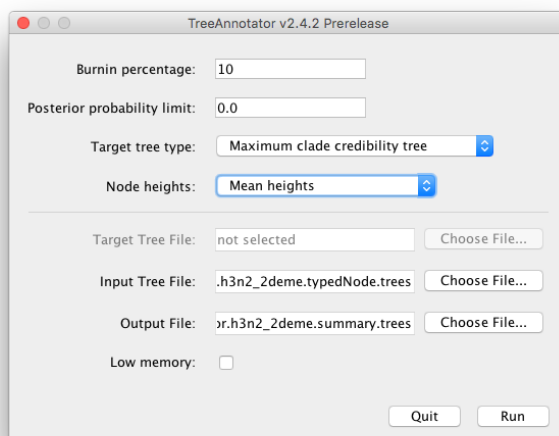


Figure 16: Use TreeAnnotator to produce a summary tree.

the “Export tree as” submenu and select “NEXUS file”. (It is important to select “NEXUS” instead of “Newick” as the Newick format does not support the annotations that MultiTypeTree uses to mark the edge types.)

4.3 Producing a summary tree using TreeAnnotator

While it is tempting to view the MAP tree shown above as the primary result of the phylogenetic side of our analysis it is very important to remember that this is only a point estimate and says nothing about the uncertainty present in the result. This is an important drawback, as we have done a full Bayesian analysis and have access to a large number of samples from the full posterior in the tree log files. The MAP tree discards almost all of this information.

We can make better use of our raw analysis results by using the TreeAnnotator program which is distributed with BEAST to analyze the `typedNode` trees which were produced by our MCMC run. To do this, simply load TreeAnnotator and select the `typedNode` tree file as the input file and `h3n2-bdmm-v2-samplingPrior.h3n2_2deme.summary.trees` as the output file. Select “Mean heights” from the “Node heights” menu and set the burn-in percentage to 10:

Pressing the “Run” button will now produce an annotated summary tree.

To visualize this tree, open IcyTree once more (maybe open it in a new browser tab), choose File->Open, then select the file `h3n2_2deme.h3n2_2deme.summary.tree` using the file selection dialog. Follow the instructions provided for the MAP tree above to colour the tree by the “type” attribute and add the legend and time axis. In addition, open the Style menu and from the “Node height error bars” sub-menu select “height_95%_HPD” to add error bars to the internal node heights. Also, open the Style menu and from the “Edge opacity” sub-menu select “type.prob”. This will cause the edges to become increasingly transparent as the posterior probability for the displayed colour decreases.

Once these style preferences have been set, you should see something similar to the following:

Here we have a full consensus tree annotated by the locations at coalescence nodes and showing node height




Figure 17: The summary tree in IcyTree.

uncertainty, with the widths of the edges representing how certain we can be of the location estimate at each point on the tree. This is a much more comprehensive summary of the phylogenetic side of our analysis.

One thing to pay attention to here is that the most probable root location is given by the summary tree to be Hong Kong (under our model which assumes that only Hong Kong and New Zealand exist). By hovering the mouse cursor over the tiny edge above the root will bring up a table in which posterior probability of the displayed root location (`type.prob`) can be seen to be approximately 90%. The analysis therefore strongly supports a Hong Kong origin over a New Zealand origin.

[Very useful final notes from Tim](#)

5 Useful Links

- Bayesian Evolutionary Analysis with BEAST 2 (Drummond and Bouckaert 2014)
- Multi-type birth-death process package (Kühnert et al. 2016)
- BEAST 2 website and documentation: <http://www.beast2.org/>



This tutorial was written by Denise Kühnert for [Taming the BEAST](#) and is licensed under a [Creative Commons Attribution 4.0 International License](#).

Version dated: June 2, 2017

Relevant References

- Drummond, AJ and RR Bouckaert. 2014. *Bayesian evolutionary analysis with BEAST 2*. Cambridge University Press,
- Kühnert, D, T Stadler, TG Vaughan, and AJ Drummond. 2016. Phylodynamics with migration: a computational framework to quantify population structure from genomic data. *Molecular Biology and Evolution*
- Vaughan, TG, D Kühnert, A Poppinga, D Welch, and AJ Drummond. 2014. Efficient bayesian inference under the structured coalescent. *Bioinformatics* 30: 2272–2279.